



الجامعة السورية الخاصة  
SYRIAN PRIVATE UNIVERSITY

Week 9

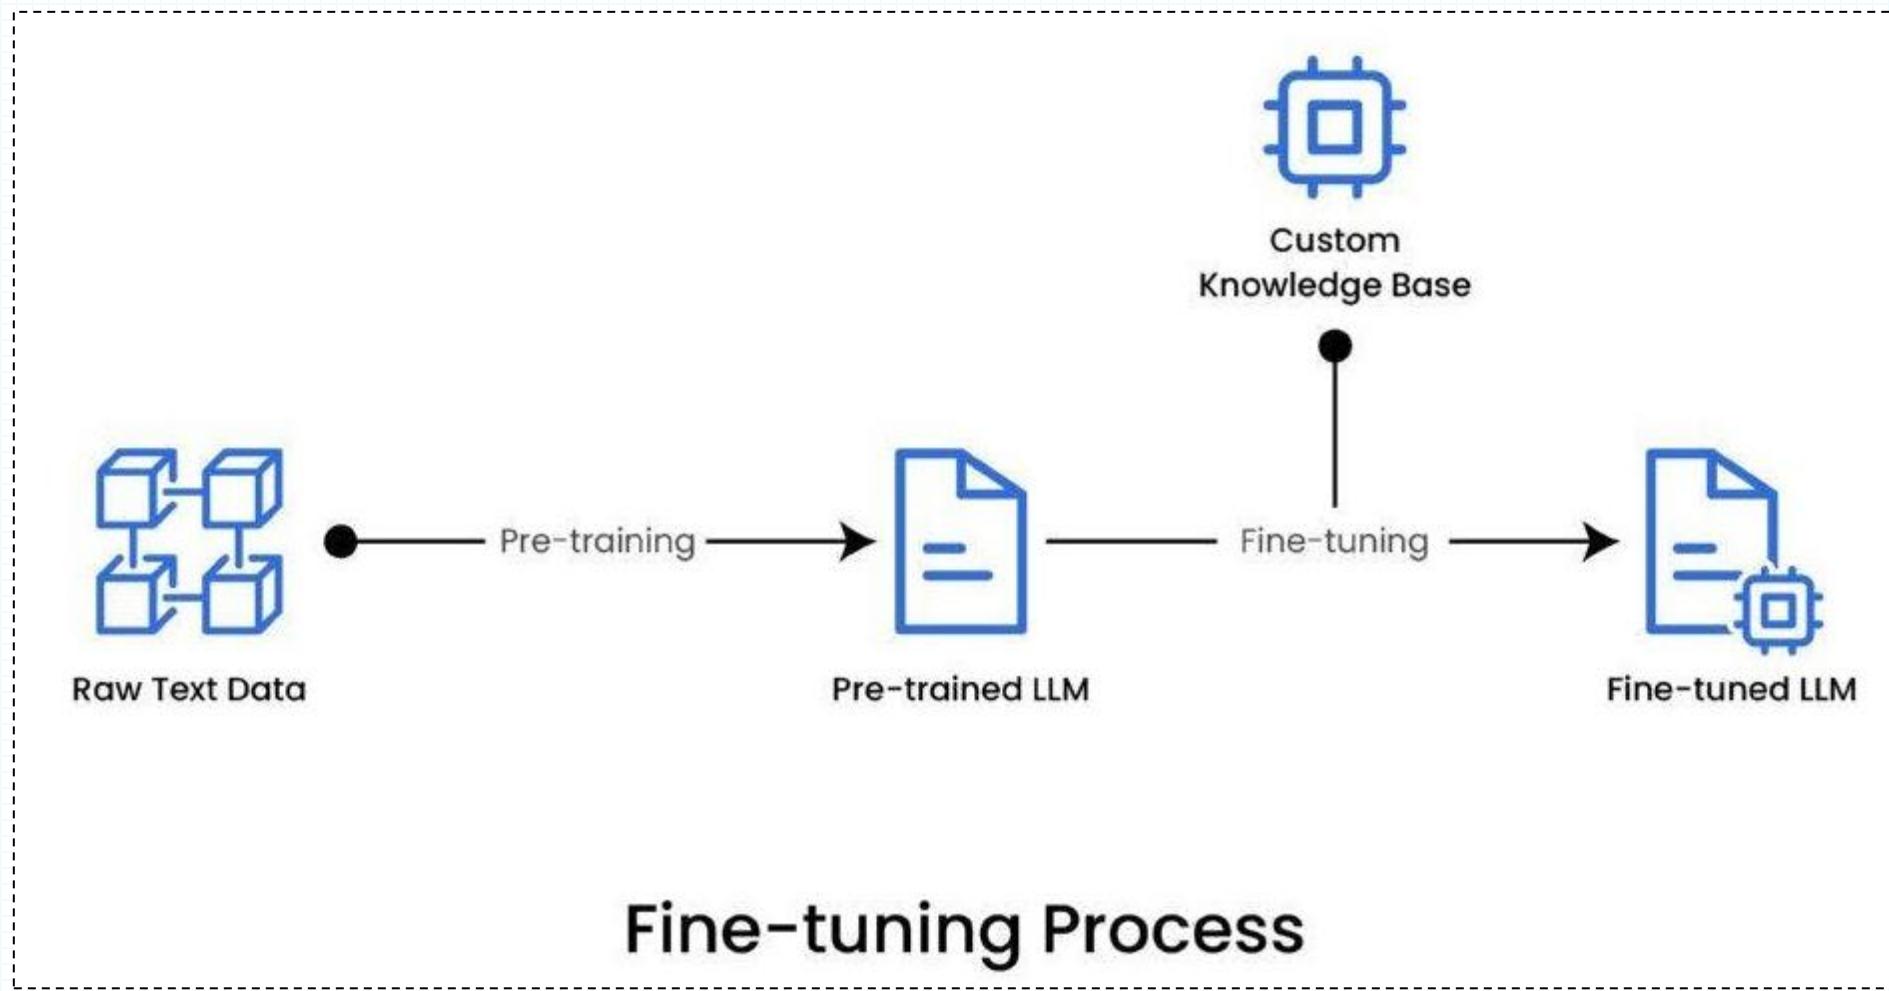
كلية الهندسة

الذكاء الصنعي العملي

# Transfer Learning: Fine-Tuning and Domain Adaptation

د. رياض سنبل

# What is LLM fine-tuning?



# What is LLM fine-tuning?

---

- Fine-tuning is the process of adjusting the parameters of a pre-trained large language model to a **specific task or domain**.
- Although pre-trained language models like GPT possess **vast language knowledge, they lack specialization** in specific areas.
- By exposing the model to task-specific examples during fine-tuning, the model can acquire a **deeper understanding of the nuances of the domain**

# Why is LLM fine-tuning important?

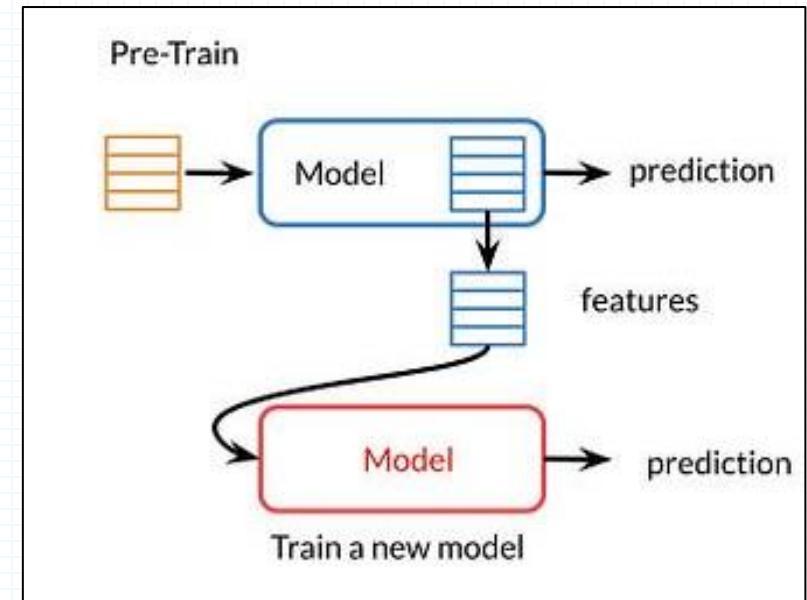
---

- **Customization:** Every domain or task has its own unique language patterns, terminologies, and contextual nuances.
- **Data compliance:** In many industries, such as healthcare, finance, and law, strict regulations govern the use and handling of sensitive information. Organizations can ensure their model adheres to data compliance standards by fine-tuning the LLM on proprietary or regulated data.
- **Limited labeled data:** Fine-tuning allows organizations to leverage pre-existing labeled data more effectively by adapting a pre-trained LLM to the available labeled dataset, maximizing its utility and performance.

# Types of LLM fine-tuning


## Feature extraction:

- It involves freezing the weights of a pre-trained model's layers and using it as a feature extractor
- Used when a smaller dataset is available, and the target domain is closely aligned with the original domain of the pre-trained model.
- Faster training, requires less computational resources, and can be more effective when the new dataset is small.



# Types of LLM fine-tuning


- **Full fine-tuning**
  - involves unfreezing some or all of the pre-trained layers and retraining them along with new layers.
  - Used when a larger dataset is available and more flexibility is needed to adapt the pre-trained model to the specific task
  - Allowing the model to learn new features that are specific to the target task
  - Can achieve higher accuracy, better adaptability to the new task, and allows the model to learn more task-specific features while preserving the general knowledge from the original dataset.



# Types of LLM fine-tuning


- **Multi-stage fine-tuning**

- A technique where a model is trained in multiple sequential stages, each stage building upon the previous one.
- The process might involve:
  - Pre-training
  - Initial Fine-Tuning: smaller dataset relevant to a specific domain e.g. medical text.
  - Domain-Specific Fine-Tuning: further fine-tuned on a dataset specifically related to a narrow sub-domain, such as clinical trial reports.




# Types of LLM fine-tuning

- **Adapter fine-tuning** is a parameter-efficient technique that allows for efficient adaptation of large language models (LLMs) to specific tasks by training small, task-specific "adapter" modules instead of updating all model parameters.
- Adapters: Adapters are small neural network layers (typically a pair of fully connected layers) inserted into the pre-trained model.





# Domain Adaptation

- Domain Adaptation (DA) is a machine learning technique to transfer knowledge from a "**source**" domain (with lots of labeled data) to a different but related "**target**" domain (often with little or no labeled data)
- Unsupervised domain adaptation (UDA) **aims to learn a predictive model for an unlabeled domain** by transferring knowledge from a separate labeled source domain.



# Unsupervised Domain Adaptation




Assumption:

The alignment will respect the label


This section contains two diagrams under the heading "Assumption:". The top diagram shows the Source domain (red triangles and blue circles) and Target domain (blue circles and red checkered squares) with a curved arrow pointing from the Source to the Target, similar to the first diagram. The bottom diagram shows the same domains with a curved arrow pointing from the Target back to the Source, suggesting a reciprocal or self-consistent alignment process.

# Unsupervised Domain Adaptation

$$L_{\text{total}}(\theta_f) = L_y(\theta_f, \theta_y) - \lambda L_d(\theta_f, \theta_d)$$



# Unsupervised Domain Adaptation



Method	<b>MNIST → USPS</b>	<b>USPS → MNIST</b>	<b>SVHN → MNIST</b>
	<b>173 → 105</b>	<b>105 → 173</b>	<b>143 → 173</b>
Source only	$0.752 \pm 0.016$	$0.571 \pm 0.017$	$0.601 \pm 0.011$
Gradient reversal	$0.771 \pm 0.018$	$0.730 \pm 0.020$	$0.739$ [16]
Domain confusion	$0.791 \pm 0.005$	$0.665 \pm 0.033$	$0.681 \pm 0.003$
CoGAN	$0.912 \pm 0.008$	$0.891 \pm 0.008$	did not converge
<b>ADDA (Ours)</b>	$0.894 \pm 0.002$	$0.901 \pm 0.008$	$0.760 \pm 0.018$