الجامعة السورية الخاصة
SYRIAN PRIVATE UNIVERSITY

Week 9 | كلية الهندسة | الذكاء الصنعي العملي

Transfer Learning: Fine-Tuning and Domain Adaptation

د. رياض سنبل

What is LLM fine-tuning?

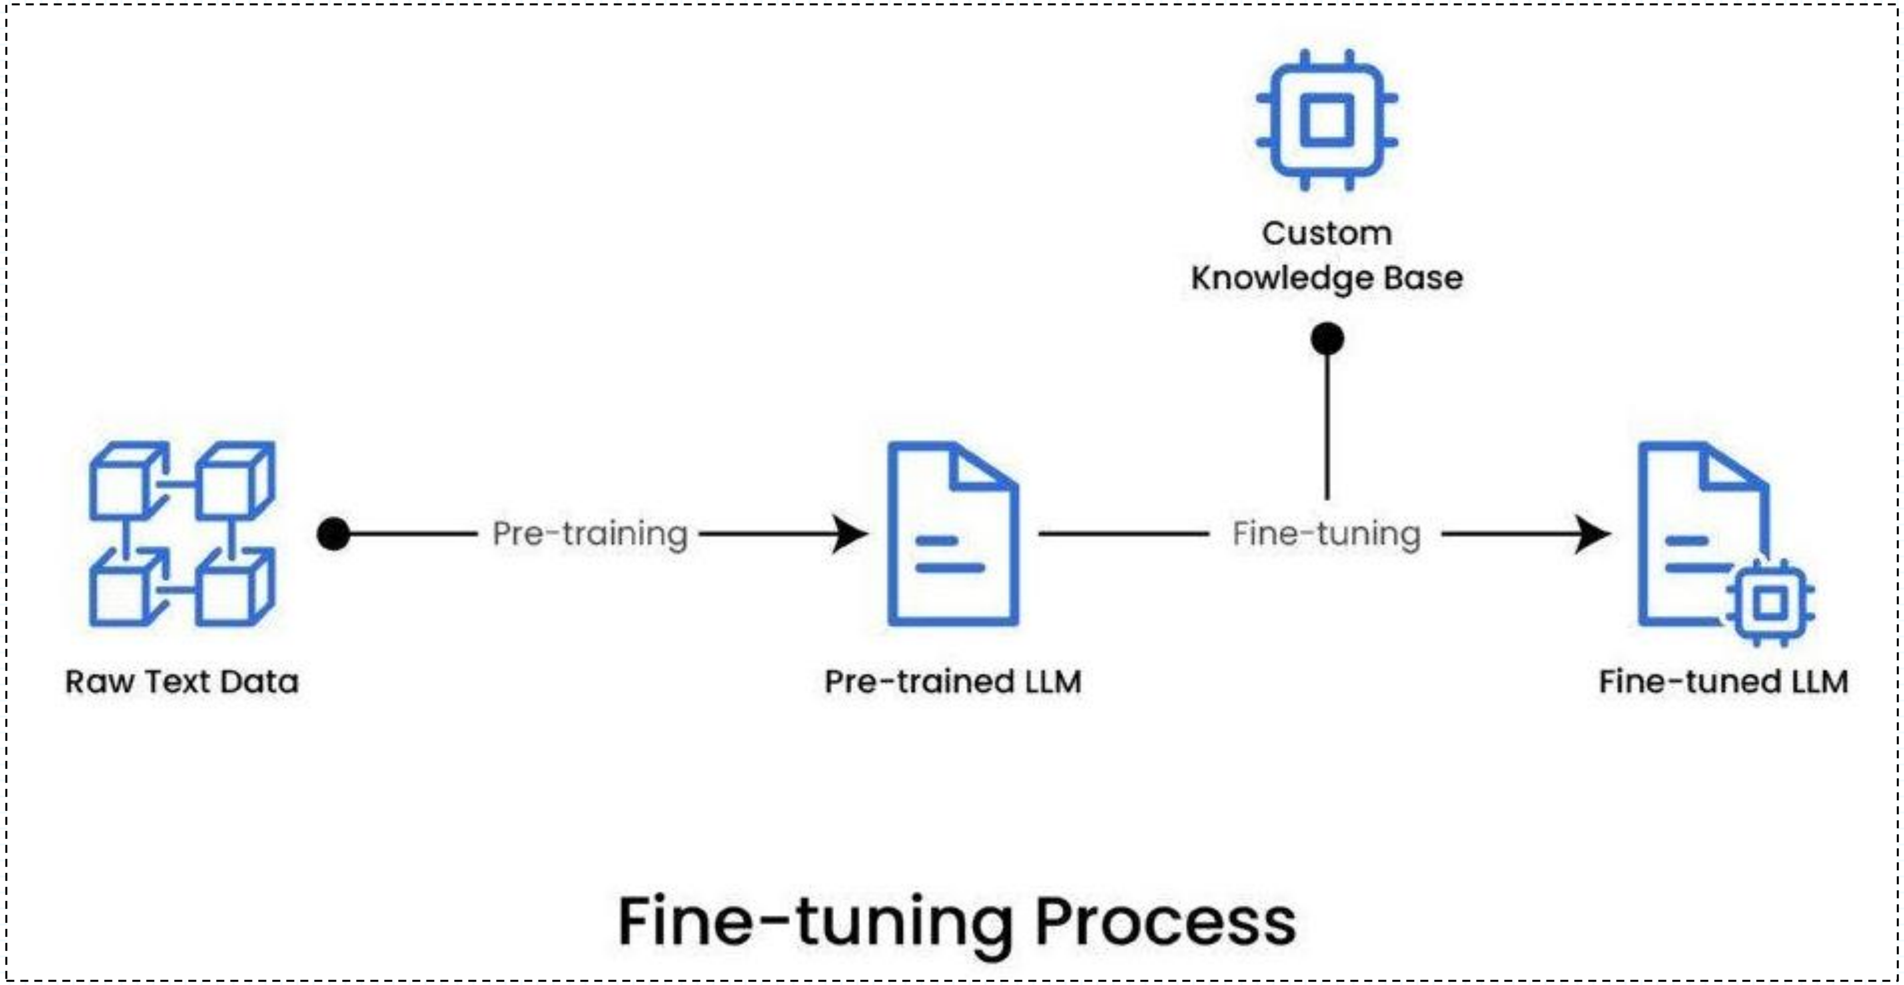

What is LLM fine-tuning?

- Fine-tuning is the process of adjusting the parameters of a pre-trained large language model to a **specific task or domain**.
- Although pre-trained language models like GPT possess **vast language knowledge**, **they lack specialization** in specific areas.
- By exposing the model to task-specific examples during fine-tuning, the model can acquire a **deeper understanding of the nuances of the domain**

Why is LLM fine-tuning important?

- **Customization**: Every domain or task has its own unique language patterns, terminologies, and contextual nuances.
- **Data compliance**: In many industries, such as healthcare, finance, and law, strict regulations govern the use and handling of sensitive information. Organizations can ensure their model adheres to data compliance standards by fine-tuning the LLM on proprietary or regulated data.
- **Limited labeled data:** Fine-tuning allows organizations to leverage pre-existing labeled data more effectively by adapting a pre-trained LLM to the available labeled dataset, maximizing its utility and performance.

Types of LLM fine-tuning

- **Feature extraction:**
  - It involves freezing the weights of a pre-trained model's layers and using it as a feature extractor
  - Used when a smaller dataset is available, and the target domain is closely aligned with the original domain of the pre-trained model.
  - Faster training, requires less computational resources, and can be more effective when the new dataset is small.

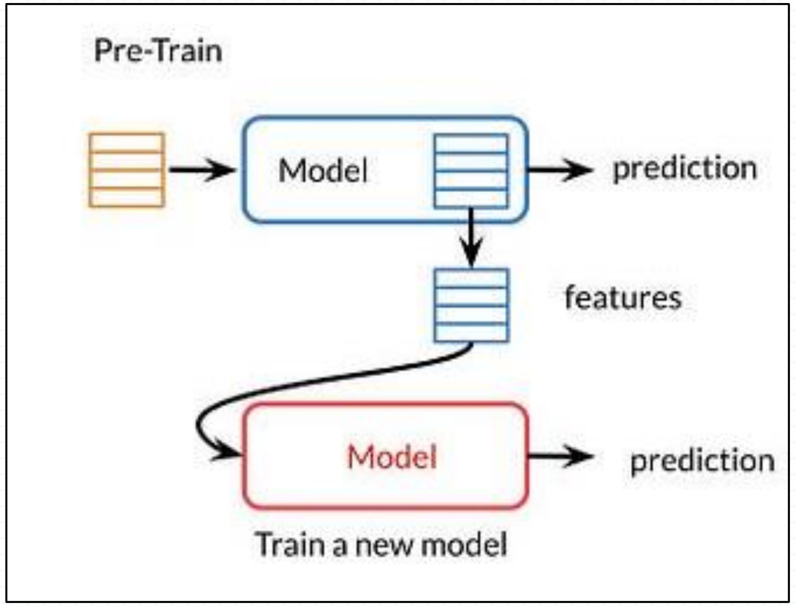

Types of LLM fine-tuning

- **Full fine-tuning**
  - involves unfreezing some or all of the pre-trained layers and retraining them along with new layers.
  - Used when a larger dataset is available and more flexibility is needed to adapt the pre-trained model to the specific task
  - Allowing the model to learn new features that are specific to the target task
  - Can achieve higher accuracy, better adaptability to the new task, and allows the model to learn more task-specific features while preserving the general knowledge from the original dataset.

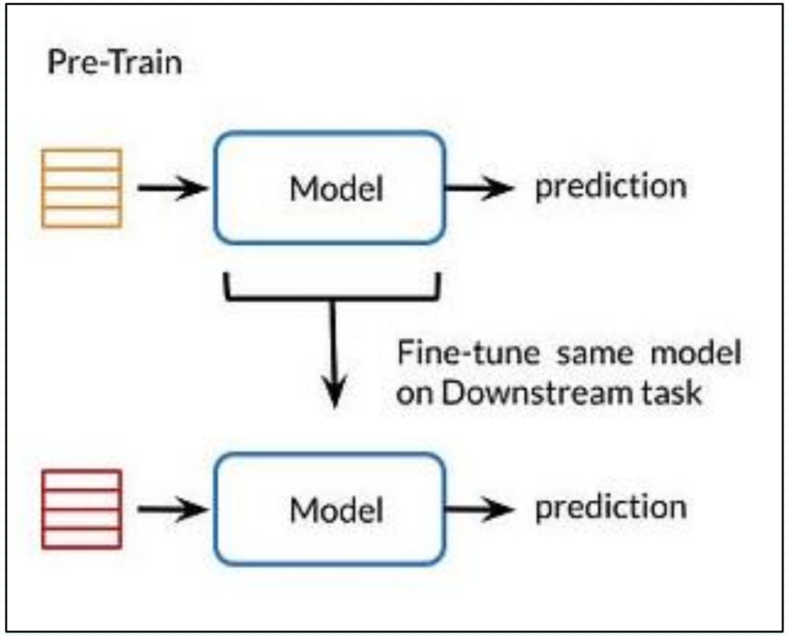

Types of LLM fine-tuning

- **Multi-stage fine-tuning**
  - A technique where a model is trained in multiple sequential stages, each stage building upon the previous one.
  - The process might involve:
    - Pre-training
    - Initial Fine-Tuning: smaller dataset relevant to a specific domain e.g. medical text.
    - Domain-Specific Fine-Tuning: further fine-tuned on a dataset specifically related to a narrow sub-domain, such as clinical trial reports.

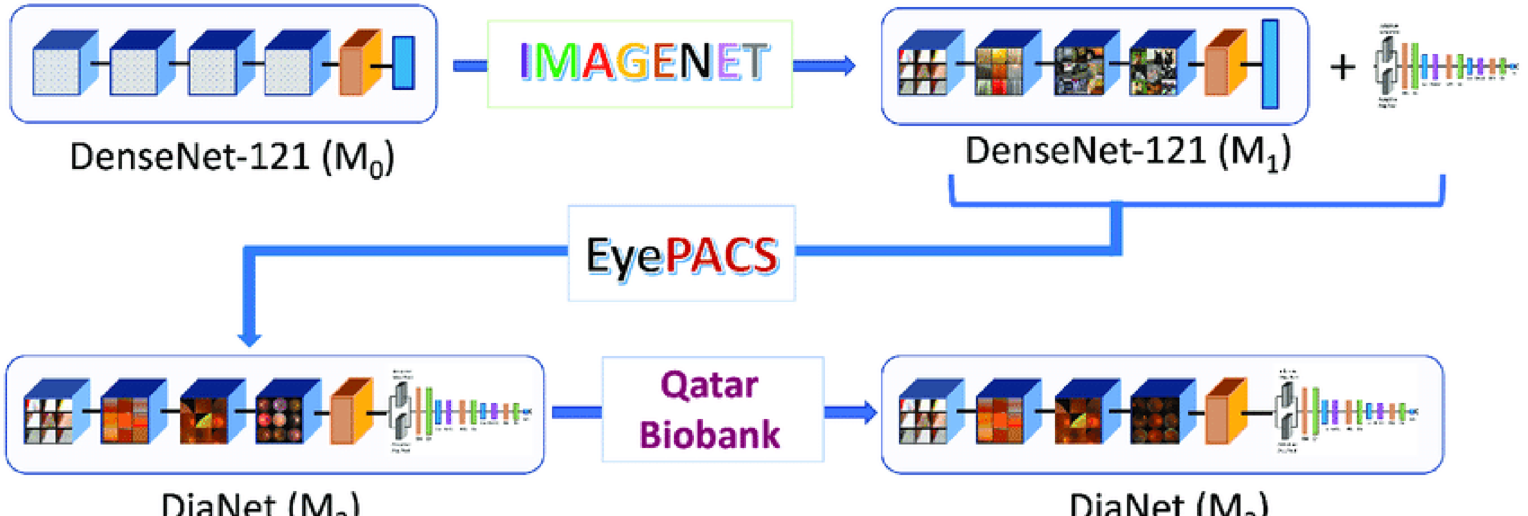

Types of LLM fine-tuning

- **Prompt Tuning**
  - Instead of manually writing a prompt ("Translate this to French: ..."), you **learn a continuous vector prompt** (a set of trainable embeddings).
  - These vectors are **prepended to the input embeddings**.

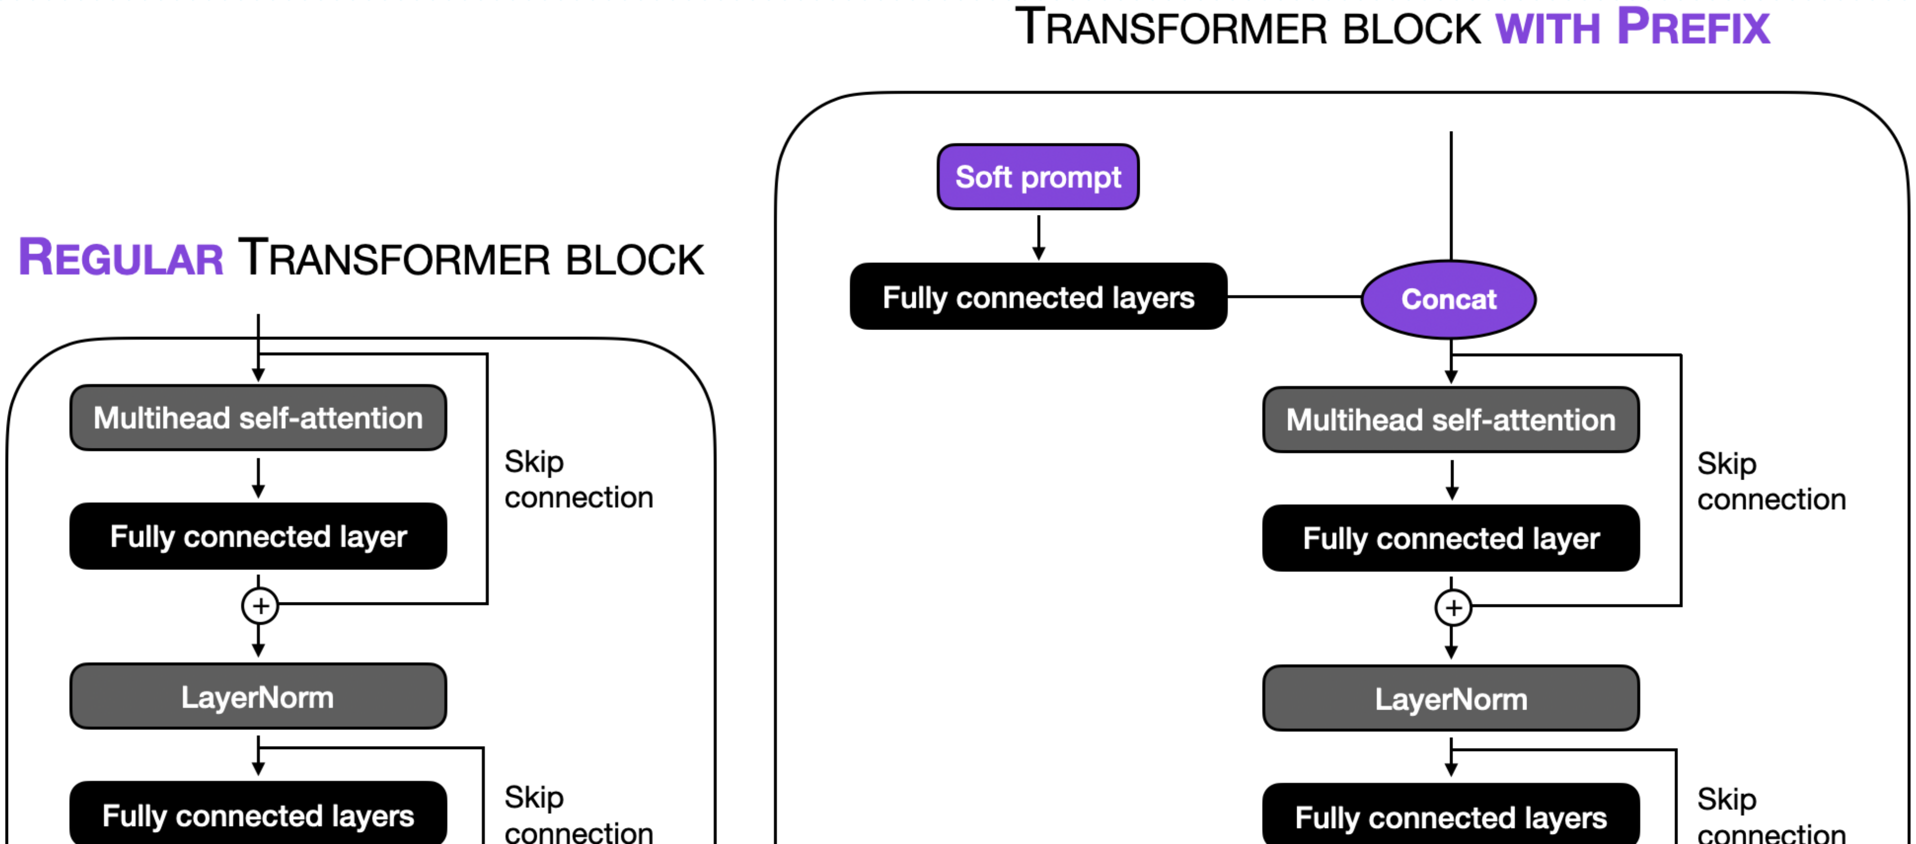

Types of LLM fine-tuning

- **Adapter fine-tuning** is a parameter-efficient technique that allows for efficient adaptation of large language models (LLMs) to specific tasks by training small, task-specific "adapter" modules instead of updating all model parameters.
- Adapters: Adapters are small neural network layers (typically a pair of fully connected layers) inserted into the pre-trained model.

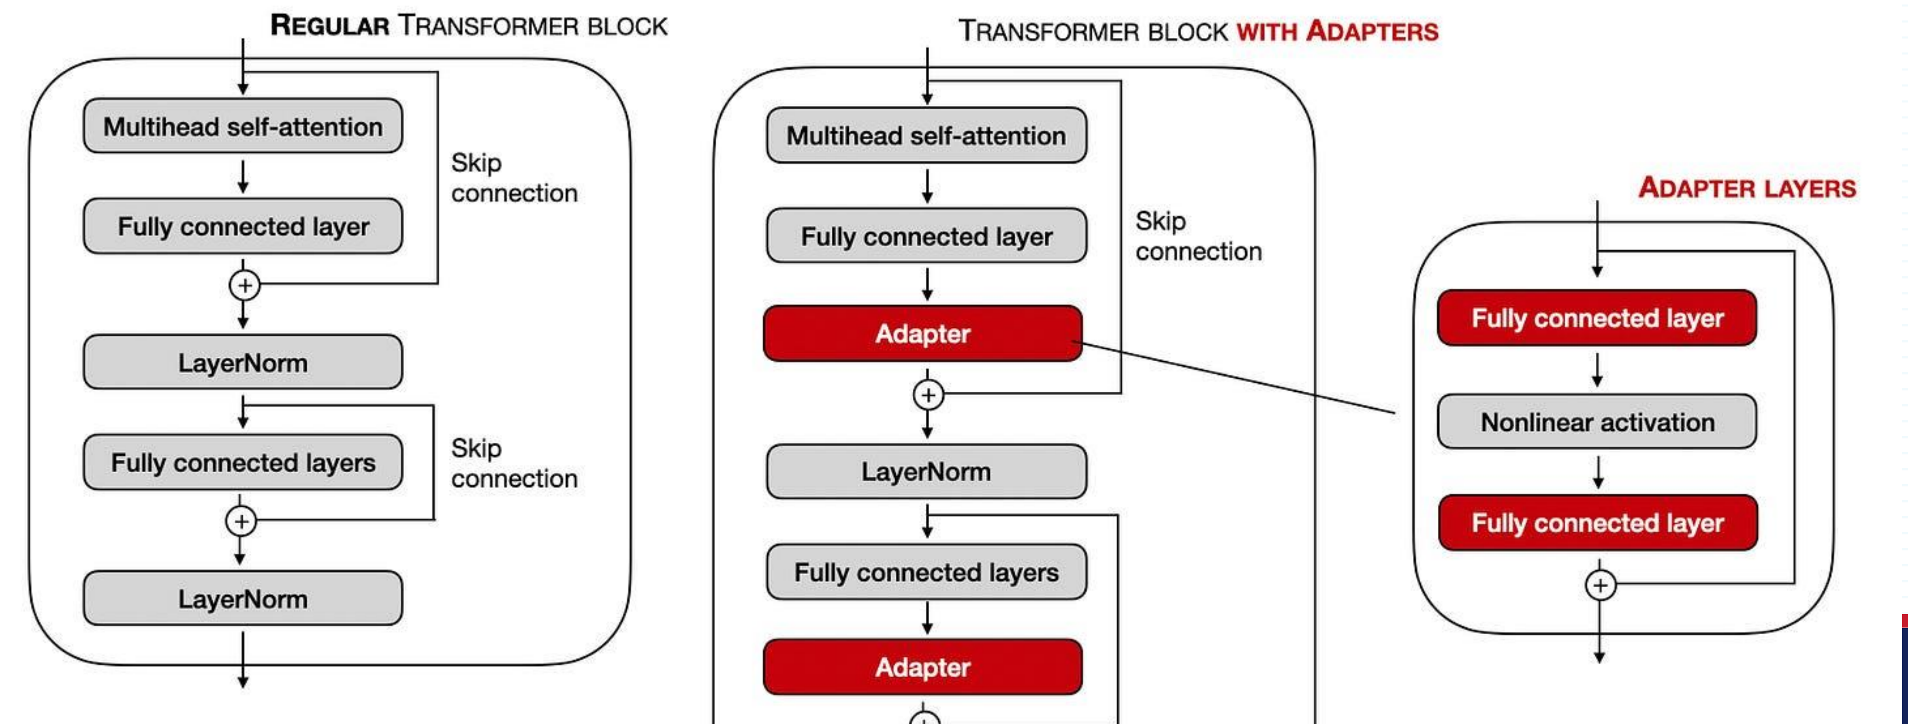

Types of LLM fine-tuning

- **LoRA (Low-Rank Adaptation)** is a parameter-efficient fine-tuning technique for large language models (LLMs) that significantly reduces the number of trainable parameters during fine-tuning.

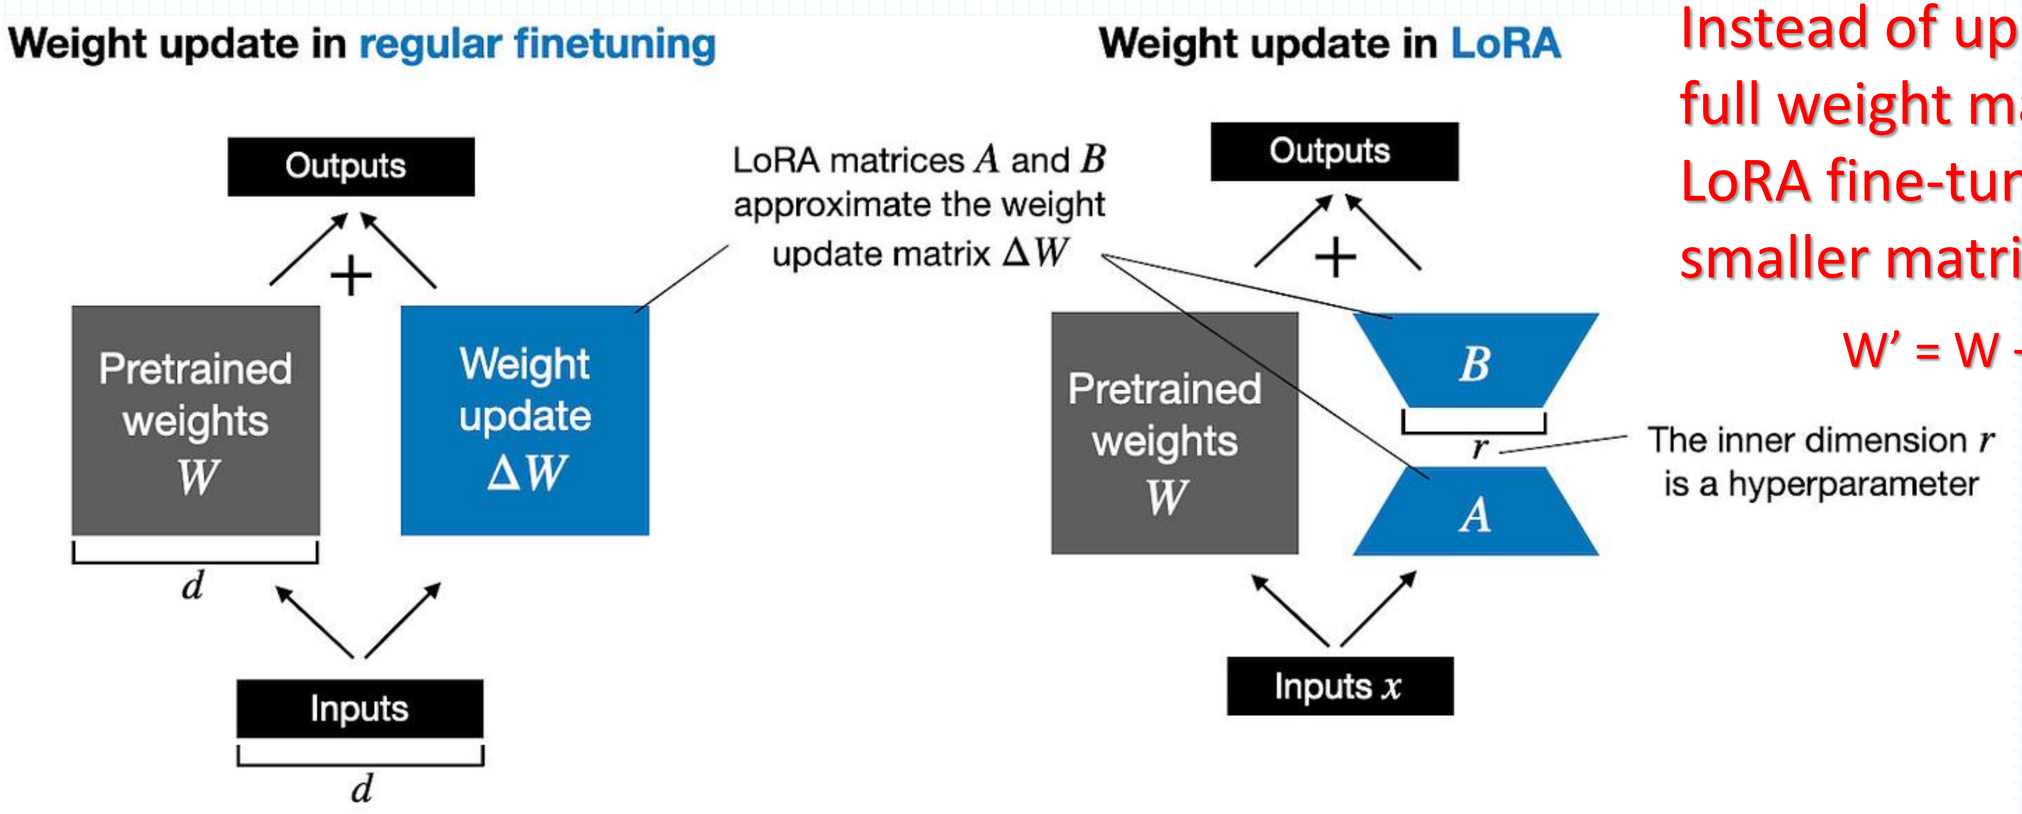

Instead of updating the full weight matrices, LoRA fine-tunes only the smaller matrices A and B

W’ = W + AB.

Few-shot Learning with LLMs

- a technique where an LLM learns to perform a new task by being presented with a few examples, typically 1 to 5, within the prompt.
- It allows the model to adapt its knowledge and capabilities to a specific task without requiring retraining or fine-tuning.
- Few-shot learning does not update the model's parameters.
- Instead, you provide a few examples in the input prompt to teach the model what kind of output you want.
- This is also called in-context learning.

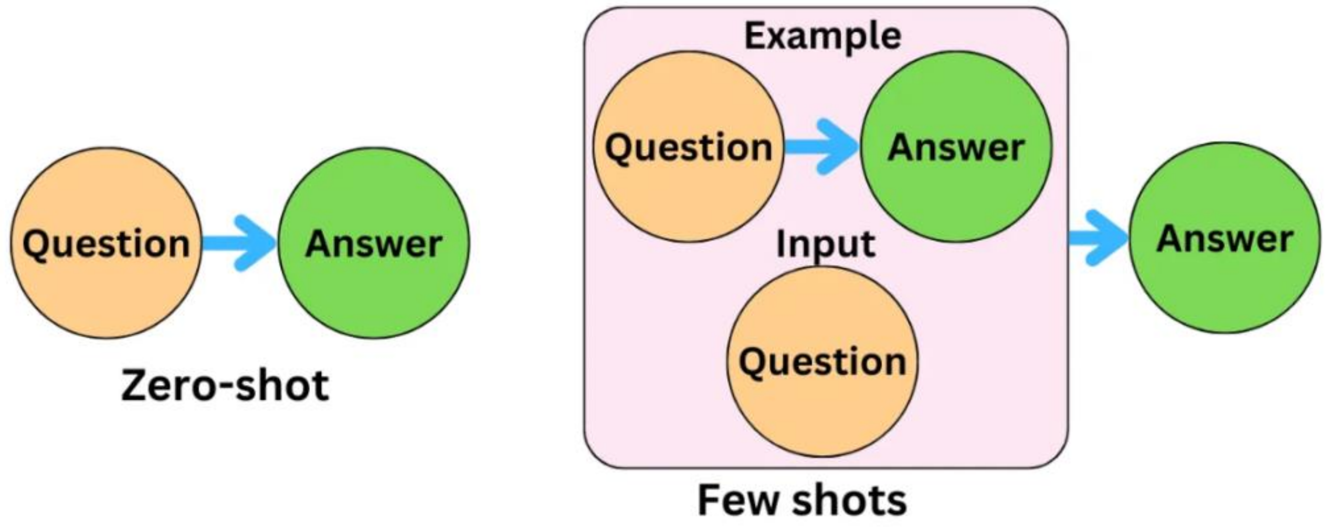

Unsupervised Domain Adaptation

Unsupervised Domain Adaptation

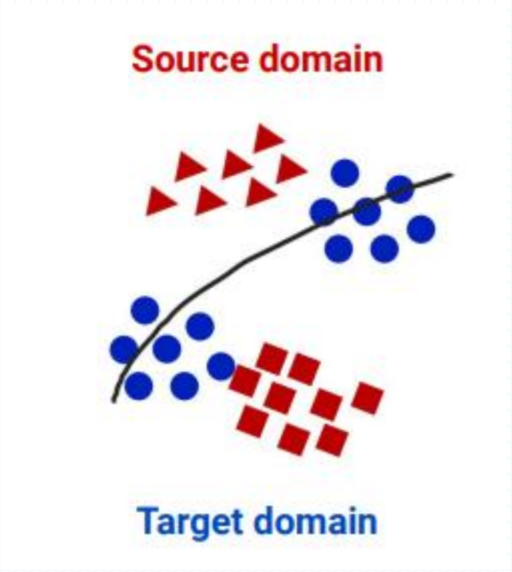

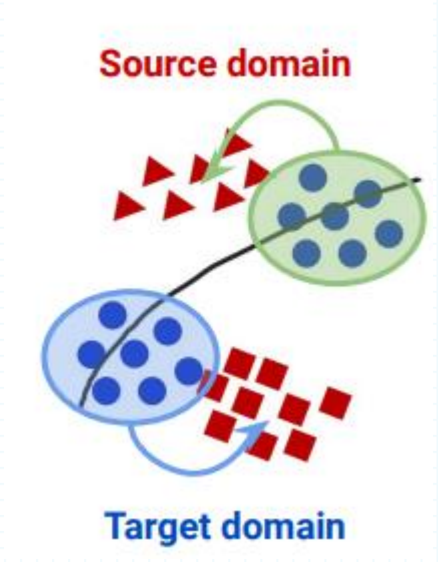

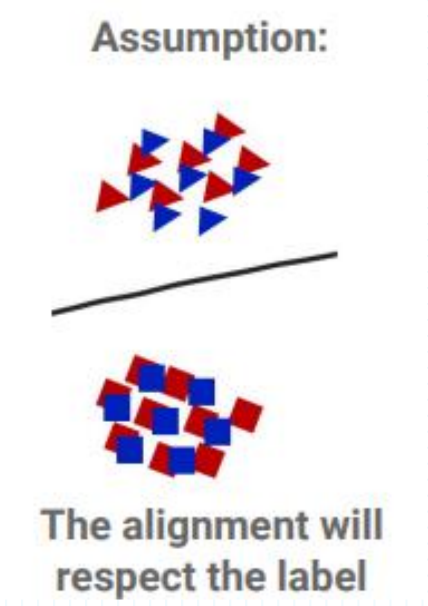

Unsupervised Domain Adaptation

$$L_{\text{total}}(\theta_f) = L_y(\theta_f, \theta_y) \; - \; \lambda \, L_d(\theta_f, \theta_d)$$

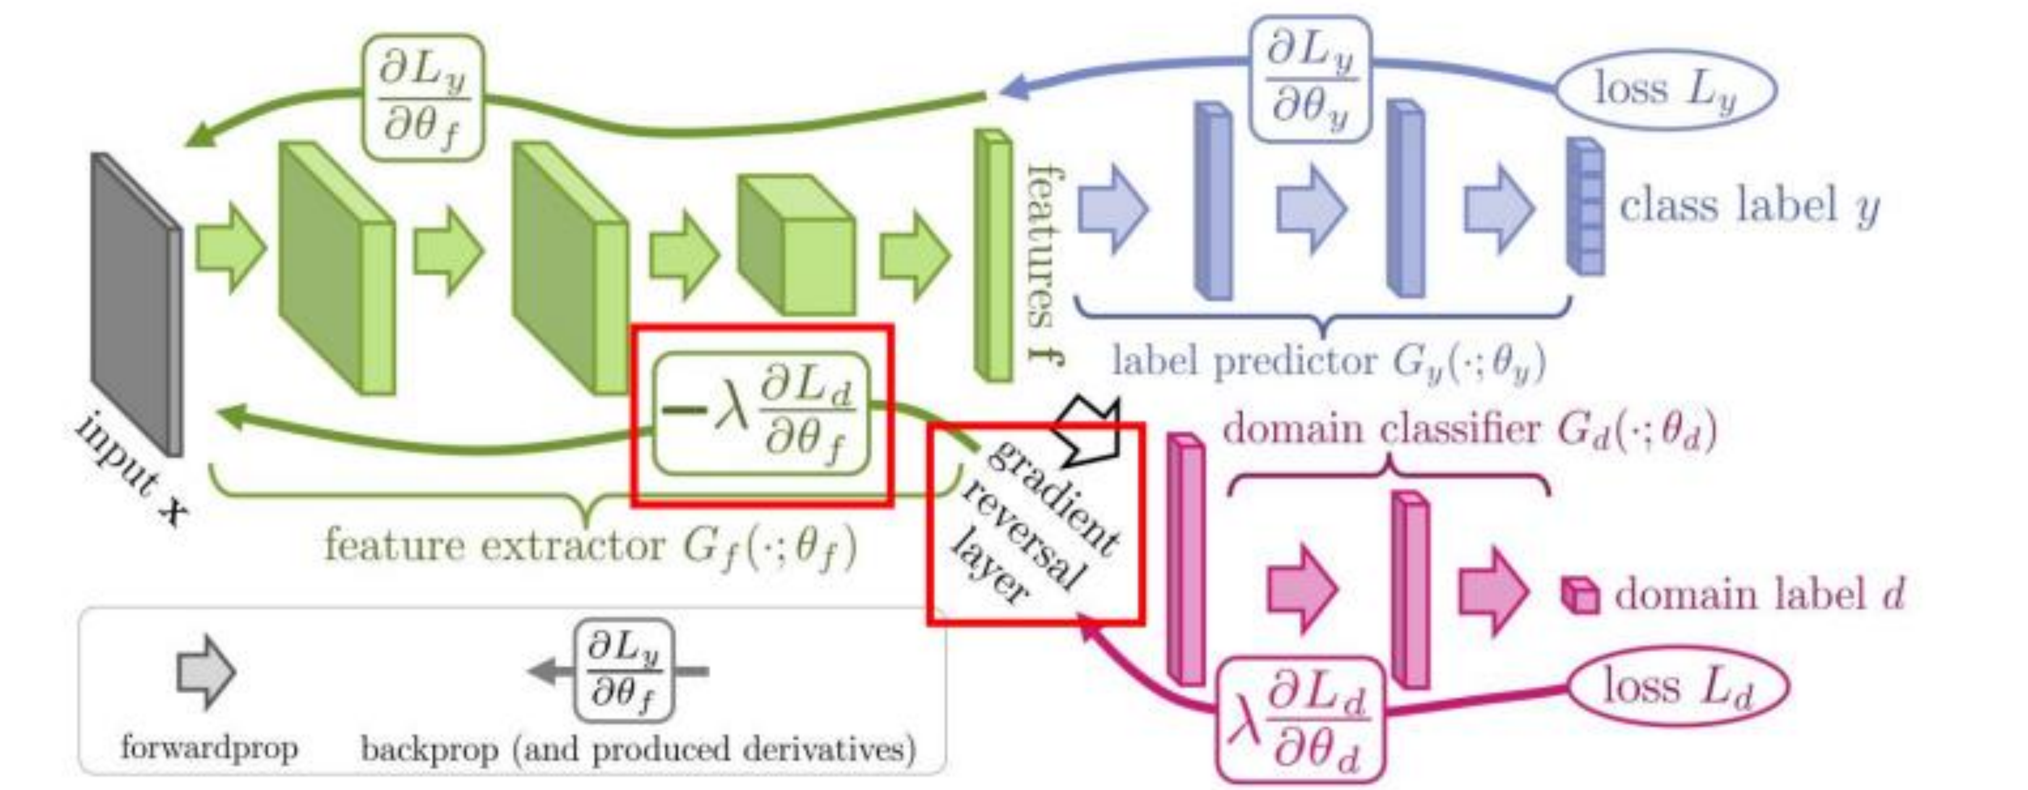

Unsupervised Domain Adaptation

Method	MNIST → USPS	USPS → MNIST	SVHN → MNIST
Source only	0.752 ± 0.016	0.571 ± 0.017	0.601 ± 0.011
Gradient reversal	0.771 ± 0.018	0.730 ± 0.020	0.739 [16]
Domain confusion	0.791 ± 0.005	0.665 ± 0.033	0.681 ± 0.003
CoGAN	0.912 ± 0.008	0.891 ± 0.008	did not converge
ADDA (Ours)	0.894 ± 0.002	0.901 ± 0.008	0.760 ± 0.018